

# CSC656 CP3 writeup

Runtime (s)							
Problem Size	blas	basic	vectorized	omp-1	omp-4	omp-16	omp-64
1024	0.00016	0.0009	0.00024	0.0009	0.00047	0.00028	0.00065
2048	0.00051	0.00358	0.0011	0.00357	0.00184	0.00167	0.00179
4096	0.00410	0.01455	0.00476	0.01454	0.00432	0.00333	0.00332
8192	0.01838	0.0583	0.01963	0.05819	0.01544	0.01304	0.01341
16384	0.06815	0.23697	0.07996	0.23310	0.06004	0.05635	0.05827

MFLOP/s							
Problem Size	blas	basic	vectorized	omp-1	omp-4	omp-16	omp-64
1024	13051.57	2321.365	8758.753	2342.598	4420.333	7601.902	3205.512
2048	16353.47	2342.743	7600.29	2351.923	4546.888	5035.478	4694.002
4096	8187.42	2306.812	7054.885	2307.514	7768.228	10080.596	10108.252
8192	7304.32	2302.219	6838.273	2306.389	8691.11	10290.996	10006.699
16384	7878.1	2265.551	6714.636	2303.132	8941.44	9528.232	9213.231

% memory bandwidth utilized							
Problem Size	blas	basic	vectorized	omp-1	omp-4	omp-16	omp-64
1024	0.02489	0.00443	0.01671	0.00447	0.00843	0.01450	0.00611
2048	0.01560	0.00223	0.00725	0.00224	0.00434	0.0048	0.00448
4096	0.0039	0.0011	0.00336	0.00011	0.0037	0.00481	0.00482
8192	0.00174	0.00055	0.00163	0.00055	0.00207	0.00245	0.00239
16384	0.00094	0.00027	0.0008	0.00027	0.00107	0.00114	0.0011

Screenshot of `dgemv_vectorized.cpp`  
turned into assembly in [godbolt.org](https://godbolt.org):

The image shows a screenshot of the Visual Studio Code editor. The left pane displays a C++ source file named 'C++ source #1'. The code defines a function `my_dgemv` that performs a vector-matrix multiplication. The function signature is `void my_dgemv(int n, double* A, double* x, double* y) {`. The implementation includes a loop for rows and a loop for columns, calculating the dot product of a row in matrix A and vector x, and storing the result in y. The right pane shows the assembly output for the same function, generated by x86-64 gcc 11.2. The assembly code is in AT&T syntax and shows the use of SSE registers (xmm) for vector operations, including `vfmovsd`, `vfmulsd`, `vfmaddsd`, and `vfmaddsd`. The assembly code is organized into blocks labeled `.L14` and `.L15`.

## Question answers:

Q1.

At  $n=16384$  between basic and vectorized, in terms of MFLOP/s, vectorized has a much better performance. The difference is  $6714.636(\text{vect}) - 2265.551(\text{basic}) = 4449.085$

In terms of memory utilization, vectorized was also better using 0.0008% vs 0.0027%. A difference of 0.0019%

Q2.

At  $n=16384$  between basic and openMP-4, in terms of MFLOP/s, omp-4 has the better performance. The difference is  $8941.44(\text{omp-4}) - 2265.551(\text{basic}) = 6675.889$

In terms of memory utilization, basic had the better performance better using 0.0027% vs 0.00107%. A difference of 0.0080%

Q3.

From 1 -> 4 threads, the speedup is:  
 $0.23310 / 0.06004 = 3.8824$

From 1 -> 16 threads, the speedup is:  
 $0.23310 / 0.05635 = 4.1366$

From 1 -> 64 threads, the speedup is:  
 $0.23310 / 0.05827 = 4.0003$